

# **Program Analysis for Software Security**

## **Lecture 1**

# Grading

Seminar activity (2-3 practical assignments + 1 oral presentation) :--**50%**

Final exam: -- **50%**

- Final Written Exam (open books): --**50%**
- **To pass the exam you have to obtain minim 5 at the final exam and the final grade is minimum 5**

# Rules

- seminar activity will be done at the group level
- groups consist of max 2 students
- final exam is individual and is an open book exam (you can have access at the lecture notes and the seminar notes)

# Course Content

Please join the Course Teams- code:

**st0nimi**

In this course we will discuss about

**Automated Program Verification**

using the **VIPER tool**.

# **Important NOTE**

**Some of the slides are taken from**

**Peter Muler, ETH Zurich**

**and**

**Christoph Matheja, DTU, Denmark**

“Program testing can be used to show the presence of bugs, but never to show their absence!”

**Edsger W. Dijkstra (1970)**

# Outline

1. Why Program Verification?
2. Course Overview
3. Starting with Viper

# How much confidence do we have in computer systems?

more confidence



Testing is insufficient

- 1994 Intel® Pentium® Floating-point Division bug
- Estimate: 1 in 9 billion floating-point divisions inaccurate
- Issue: missing entries in the lookup table
- Recall losses: \$475 million (> 5 billion DKK in 2019)
- Bug was detected during experiments on number theory

extensive testing

no confidence



# How much confidence do we have in computer systems?

more confidence



extensive testing

no confidence

OpenJDK's `java.util.Collection.sort()` is broken:  
The good, the bad and the worst case\*

Stijn de Gouw<sup>1,2</sup>, Jurriaan Rot<sup>3,1</sup>, Frank S. de Boer<sup>1,3</sup>, Richard Bubel<sup>4</sup>, and  
Reiner Hähnle<sup>4</sup>

- TimSort: default sorting algorithm in OpenJDK and Android SDK
- Certain large arrays ( $\geq 67M$ ) lead to index-out-of-bounds errors
- Multiple attempts to fix related errors were ineffective

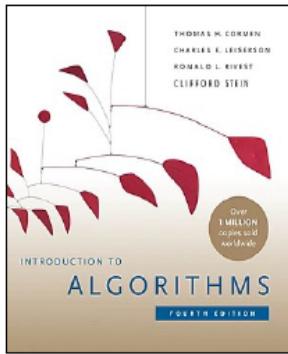
Program testing can be very effective to show the presence of bugs, but it is **hopelessly inadequate** for showing their absence.



Edsger W. Dijkstra

# How much confidence do we have in computer systems?

more confidence



correctness arguments

extensive testing

no confidence

The [only effective way](#) to raise the confidence level of a program is to give a convincing [proof of its correctness](#).



Edsger W. Dijkstra

PARTITION( $A, p, r$ )

```
1  $x = A[r]$ 
2  $i = p - 1$ 
3 for  $j = p$  to  $r - 1$ 
4   if  $A[j] \leq x$ 
5      $i = i + 1$ 
6     exchange  $A[i]$  with  $A[j]$ 
7 exchange  $A[i + 1]$  with  $A[r]$ 
8 return  $i + 1$ 
```

At the beginning of each loop iteration:

1. If  $p \leq k \leq i$ , then  $A[k] \leq x$ .
2. If  $i + 1 \leq k \leq j - 1$ , then  $A[k] > x$ .
3. If  $k = r$ , then  $A[k] = x$ .

credits: Cormen et al., Introduction to Algorithms, 2009

# Textbook-style correctness arguments are insufficient

- 2006 implementation of binary search in `java.util.Arrays`
  - Problem: `mid` might overflow!
  - This bug was part of the Java standard library for approximately **nine years**
- **Faithful implementation** of algorithm from *Programming Pearls, Bentley, 1986*
  - **Correctness argument:** “sets `mid` to the average of `low` and `high`, truncated down to the nearest integer”
  - It was inconceivable at the time to have arrays of length  $2^{30}$  or greater
- Same issue arises for other divide-and-conquer algorithms

```
public static int binarySearch(  
    int[] a, int key) {  
    int low = 0;  
    int high = a.length - 1;  
  
    while (low <= high) {  
        int mid = (low + high) / 2;  
        int midVal = a[mid];  
  
        if (midVal < key)  
            low = mid + 1;  
        else if (midVal > key)  
            high = mid - 1;  
        else  
            return mid; // key found  
    }  
    return -(low + 1); // key not found  
}
```

# How much confidence do we have in computer systems?

more confidence



rigorous proofs

correctness arguments

extensive testing

no confidence

The **only effective way** to raise the confidence level of a program is to give a convincing **proof of its correctness**.



Edsger W. Dijkstra,  
ACM Turing Lecture 1972

```
{ true }
if (x < 0) {
    { -x + 1 = |-x| + 1 }
    x := -x
    { x + 1 = |x| + 1 }
}
{ x + 1 = |x| + 1 }
y := x + 1
{ y = |x| + 1 }
```

# Handwritten proofs are insufficient

## Chord: A Scalable Peer-to-peer Lookup Service for Internet Applications

Ion Stoica, Robert Morris, David Karger, M. Frans Kaashoek, Hari Balakrishnan  
MIT Laboratory for Computer Science  
[chord@lcs.mit.edu](mailto:chord@lcs.mit.edu)  
<http://pdos.lcs.mit.edu/chord/>

- Chord is a distributed hash table developed at MIT
  - Three features that distinguish Chord from many other peer-to-peer lookup protocols are its simplicity, provable correctness, and provable performance.
- None of the seven properties claimed invariant of the original version is actually an invariant

*“Beware of bugs in the above code;  
I have only proved it correct, not tried it.”*

Donald Knuth, 1977

# How much confidence do we have in computer systems?

more confidence



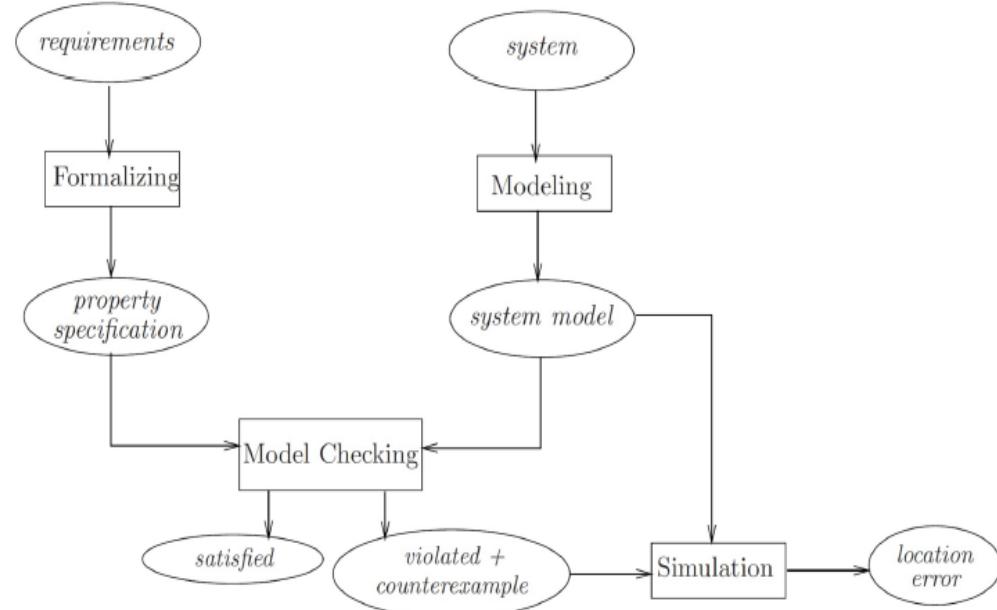
verified models

rigorous proofs

correctness arguments

extensive testing

no confidence



if **model = system**, we attempt to generate  
a rigorous proof by exhaustive testing

credits: Baier & Katoen, Principles of Model Checking, 2008

# Verification of system models is insufficient

*“Any verification using model-based techniques is only as good as the model of the system.” – Baier & Katoen 2008*

- No guarantees for the actual system
- Model checking suffers from the **state-space explosion** problem
- Not generally applicable to **infinite-state** systems
  - Ill-suited if data ranges over infinite domains (e.g., dynamic memory allocation)
  - Does not work well for systems with an arbitrary number of components

# How much confidence do we have in computer systems?

more confidence

machine-checked proofs

← our focus: deductive verification tools

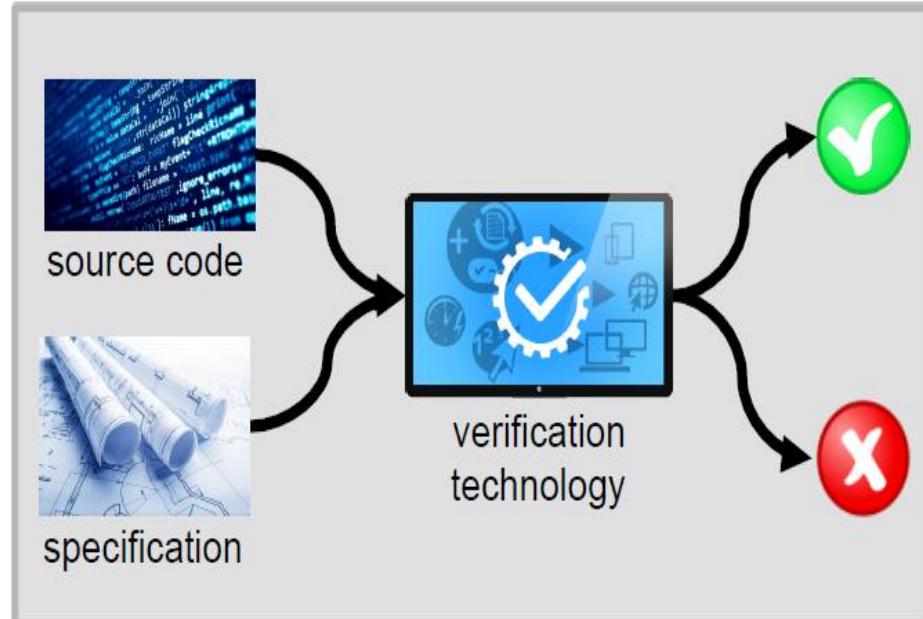
verified models

(handwritten) rigorous proofs

correctness arguments

extensive testing

no confidence



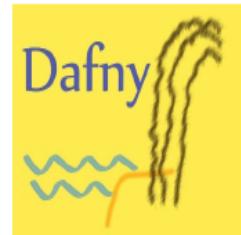
# Interactive verification

- Successful large-scale applications of interactive verification tools:
  - CompCert: a formally verified C compiler (2008)
  - seL4: a formally verified high-performance operating system microkernel (2009)
  - EveryCrypt: a formally verified crypto library (2020)
- Strengths
  - Expressive foundation (higher-order logic)
  - Can handle complex systems and properties
- Weaknesses
  - Requires **expert knowledge**
  - Very **labor-intensive** (CompCert required more than 6 person years)



# Automated (or auto-active) Verification

- Idea: “use verification like compilation”
  - Specifications take the form of source code annotations
  - Analogies: TypeScript, Rust ownership & traits, Python type hints
- Strengths:
  - Substantially less effort than interactive verification
  - Integrates into existing development processes
  - More annotations → more correctness guarantees
- Weaknesses:
  - Less expressive than interactive verification
  - May produce false positives (due to undecidability)
  - Still requires effort and expertise



P\*rust-\*i



# Outline

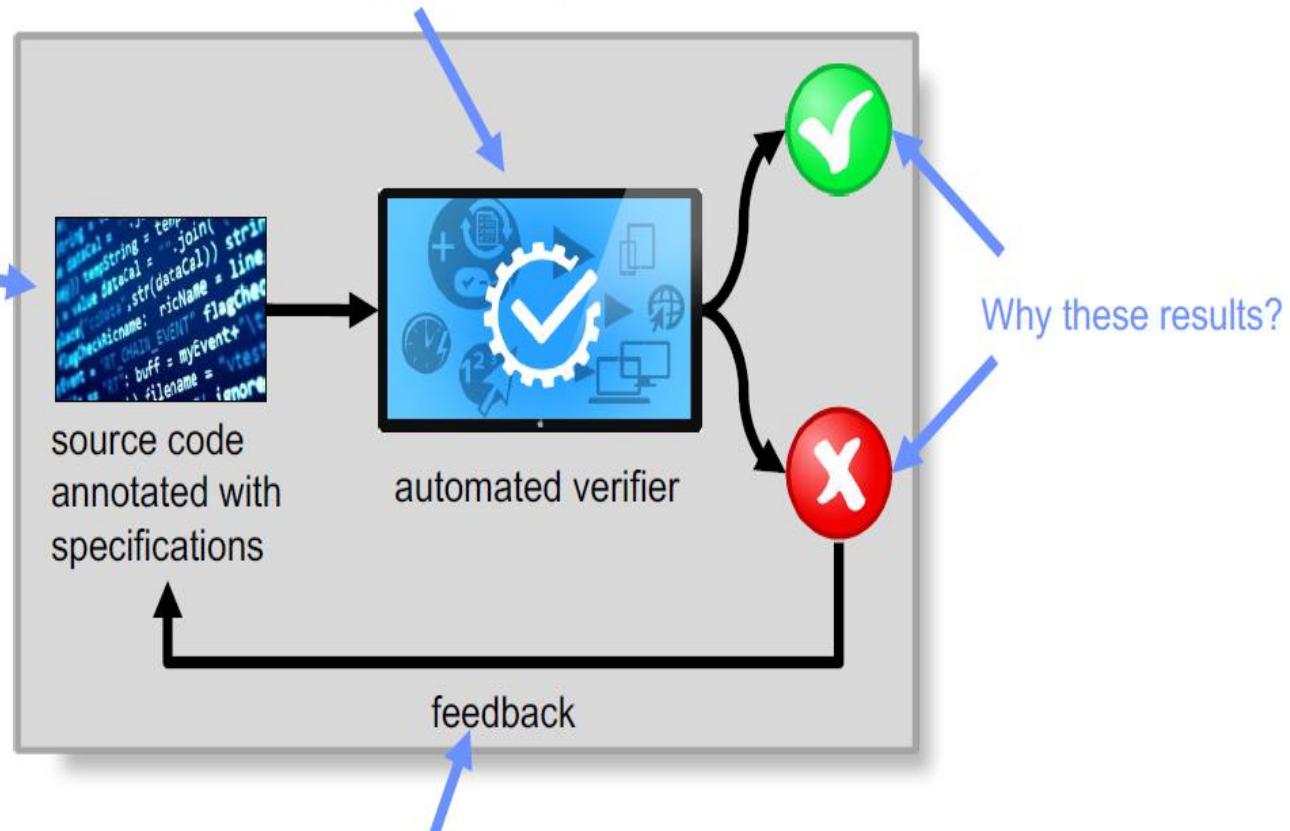
1. Why Program Verification?
2. Course Overview
3. Starting with Viper

# Course objectives

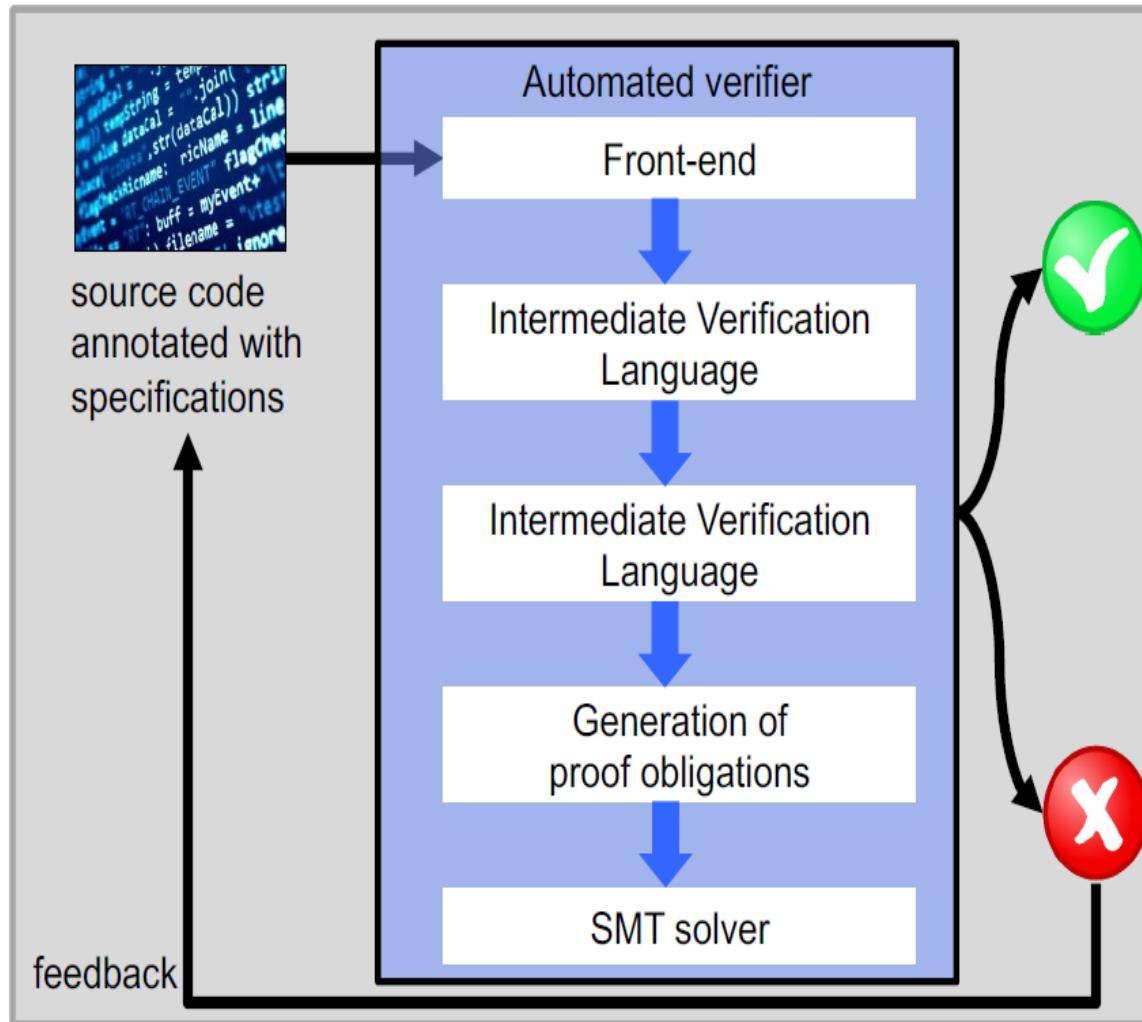
How does this work?

How do we apply and implement this?

How to specify properties?  
How to write good specifications?

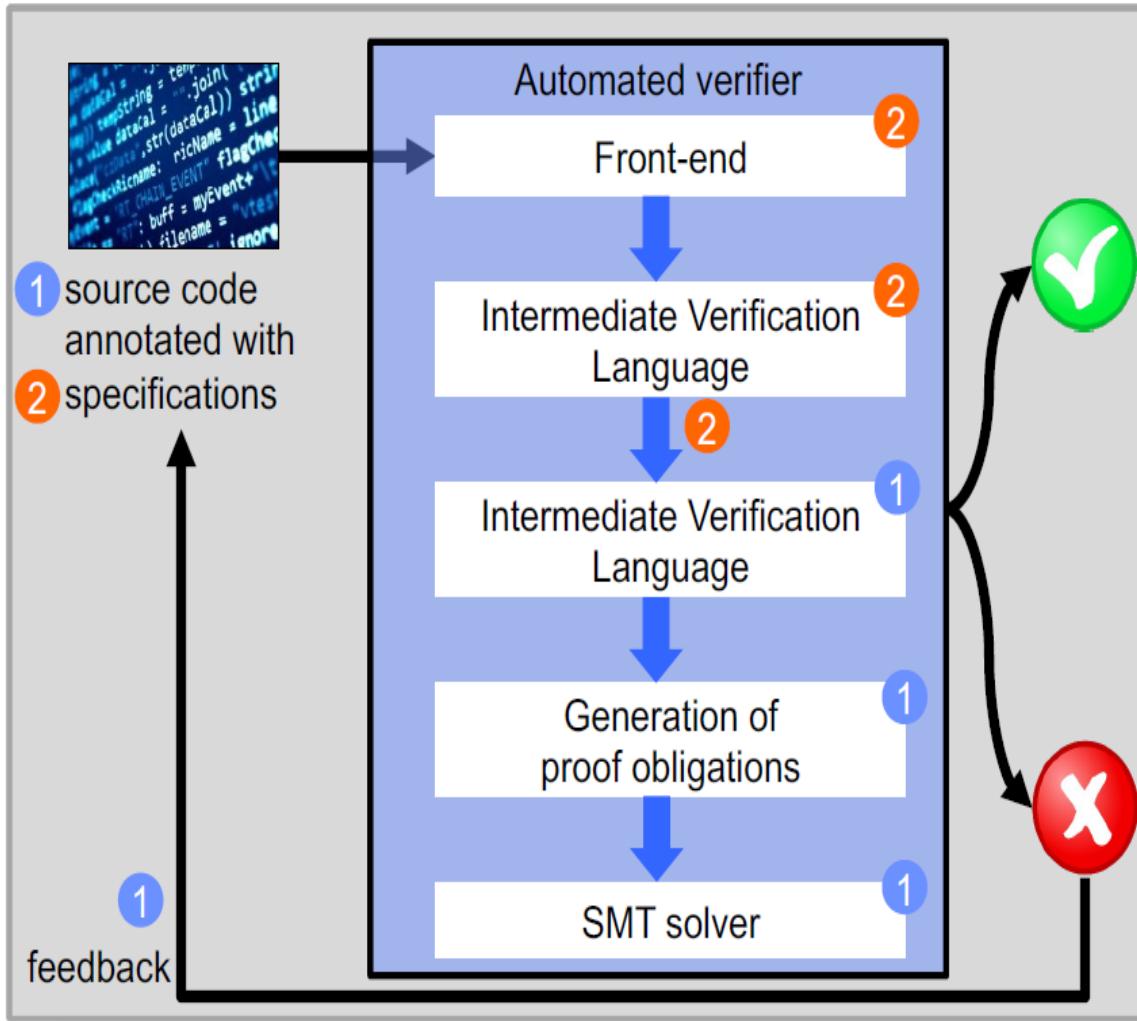


# Architecture of automated program verifiers



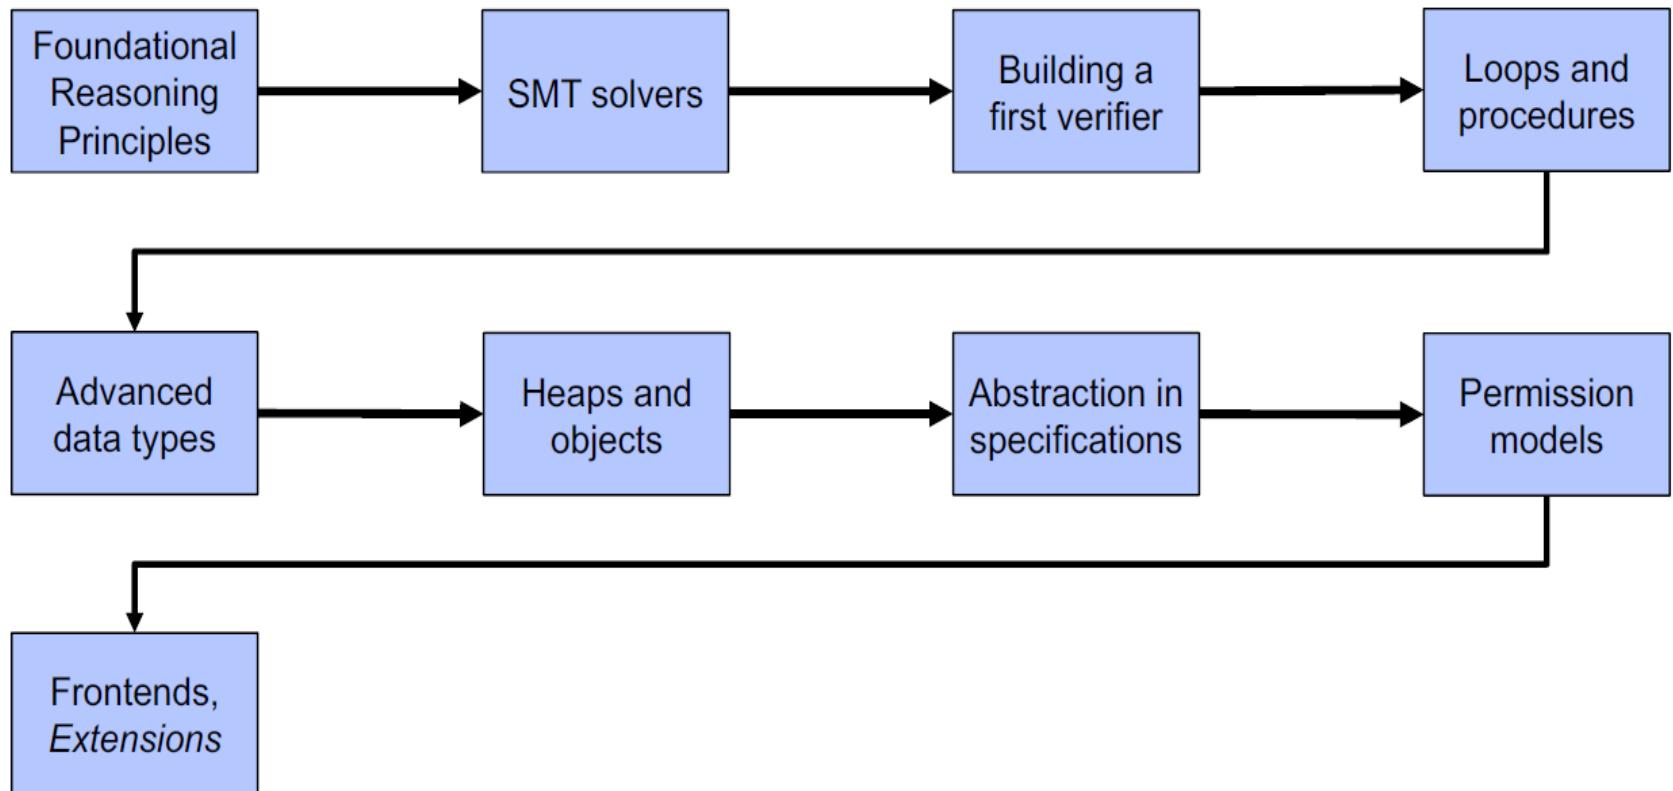
- Automated verifiers are often implemented as a tool stack
- Stepwise [compilation](#) of programs into logical formulas (and back for error reporting)
- Each transformation deals with one verification problem
- Requirements:
  - reasoning principles
  - verification methodologies
  - engineering practices

# Roadmap



1. We learn how to build and use a verification tool for a small programming language
  - Core reasoning principles
  - Generation of proof obligations
  - Working with SMT solvers
  - Error reporting
2. We extend the language by advanced features
  - Verification challenges
  - Advanced reasoning and specification principles
  - Automation via encoding to lower levels

# Tentative course outline

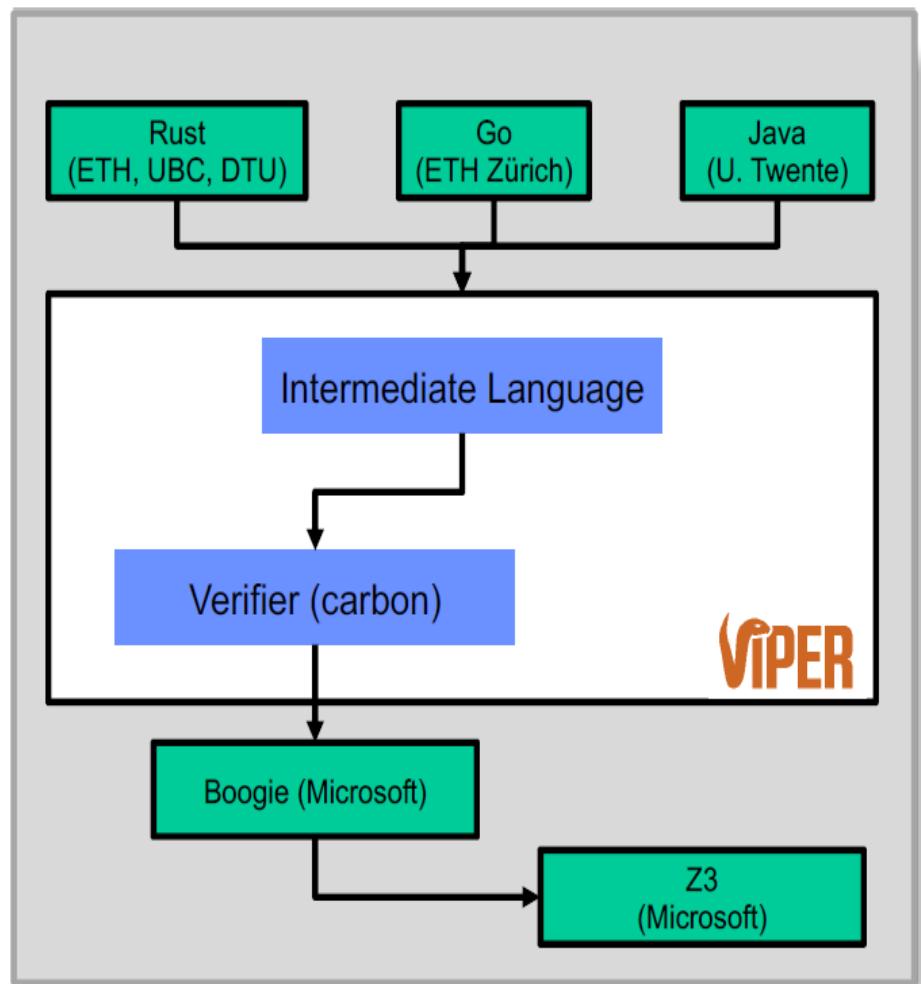


# Outline

1. Why Program Verification?
2. Course Overview
3. Starting with Viper

# The Viper Verification Framework

- Viper language
  - Models verification problems
  - Some statements are not executable
- Two verification backends
  - Carbon (close to what you will build)
  - Silicon
- **For now:** Programming language with a built-in verifier
- **Later:** Automate new methodologies



## Viper methods

```
method triple(x: Int) returns (r: Int)
{
    r := 3 * x
}
```



# Viper methods

```
method triple(x: Int, flag: Bool)
    returns (r: Int)
{
    if (flag) {
        r := 3 * x
    } else {
        var y: Int
        y := x + x
        r := x + y
    }
}
```

all Viper statements are put in **methods**

read-only **input** parameters

read-write **output** parameters

no explicit return statements

local variable declaration



# Assertions

```
method triple(x: Int, flag: Bool)
  returns (r: Int)
{
  if (flag) {
    r := 3 * x
    assert r > 0
  } else {
    var y: Int
    y := x + x
    r := x + y
    assert r == 3 * x
  }
}
```

- **assert expr** tests if expr evaluates to true
  - Yes: no effect
  - No: runtime error
- Testing: no assertion error for *chosen* inputs
- Verification: no assertion error for *all* inputs

Which assertions hold?

# Assertions

```
method triple(x: Int, flag: Bool)
    returns (r: Int)
{
    if (flag) {
        r := 3 * x
        assert r > 0 
    } else {
        var y: Int
        y := x + x
        r := x + y
        assert r == 3 * x 
    }
}
```

# Postconditions

```
method triple(x: Int) returns (r: Int)
    ensures r == 3 * x
{
    var y: Int
    y := x + x
    r := x + y
}

method client() {
    var z: Int
    z := triple(7)
    assert z == 21
}
```

- Postconditions specify how returned outputs are related to inputs
  - Default: true

# Postconditions

```
method triple(x: Int) returns (r: Int)
    ensures r == 3 * x
{
    var y: Int
    y := x + x
    r := x + y
}
check: r == 3 * x

method client() {
    var z: Int
    z := triple(7)
    assert z == 21
}
learn: z == 3 * 7
```

- Postconditions specify how returned outputs are related to inputs
  - Default: true
- Checked against implementation for all possible parameters
- Guaranteed to hold after **method calls** for supplied parameters

# Alternative Implementation

```
method triple(x: Int) returns (r: Int)
  ensures r == 3 * x
{
  r := x / 2
  r := 6 * r
}
```

x = 7  
x = 3  
x = 18 

```
method client() {
  var z: Int
  z := triple(7)
  assert z == 21
}
```

- Some implementations do not work for arbitrary inputs
- A **precondition** filters out undesirable inputs

# Preconditions

```
method triple(x: Int) returns (r: Int)
    requires x % 2 == 0
    ensures r == 3 * x
{
    r := x / 2
    r := 6 * r
}
```

```
method client() {
    var z: Int
    z := triple(7)
    assert z == 21
}
```

- **Preconditions** specify on what inputs a method can be called
  - Default: true

# Preconditions

```
method triple(x: Int) returns (r: Int)
    requires x % 2 == 0
    ensures r == 3 * x
{
    r := x / 2
    r := 6 * r
}

method client() {
    var z: Int
    z := triple(7)
    assert z == 21
}
```

 r == 3 \* x for even x

 7 % 2 == 1

- **Preconditions** specify on what inputs a method can be called
  - Default: true
- *Guaranteed* at the beginning of method implementation
- *Checked* before **method calls** for supplied parameters

# Contracts

A method **contract** consist of the method's

- name,
- input and output parameters, and
- pre- and postconditions.

Contracts must be upheld by method calls and implementations.

```
method triple(x: Int) returns (r: Int)
    requires x % 2 == 0
    ensures r == 3 * x
{
    // implementation
    r := x / 2
    r := 6 * r
}
```



```
method client()
{
    triple(7) X
    // violates precond.
}
```

# Underspecification

```
method triple(x: Int) returns (r: Int)
    requires x > 3
    ensures r > x
{
    r := 3 * x
}
```



- Implementation details are often irrelevant
- Contracts may
  - require more than an implementation needs
  - ensure less than an implementation gives

Give another contract implementation.

# Underspecification

```
method triple(x: Int) returns (r: Int)
    requires x > 3
    ensures r > x 
{
    r := 3 * x
}
```

```
method triple(x: Int) returns (r: Int)
    requires x > 3
    ensures r > x 
{
    r := x + 1
}
```

- Implementation details are often irrelevant
- Contracts may
  - require more than an implementation needs
  - ensure less than an implementation gives

Give another contract implementation.

# Verifying Method Calls

```
method triple(x: Int) returns (r: Int)
    requires x > 0
    ensures r > x
{
    r := 3 * x
}
```

```
method client() {
    var z: Int
    z := triple(7)
    assert z > 5
    assert z == 21
}
```



What is happening here?

# Verifying Method Calls

```
method triple(x: Int) returns (r: Int)
    requires x > 0
    ensures r > x
{
    r := 3 * x
}
```

```
method client() {
    var z: Int
    z := triple(7)
    assert z > 5
    assert z == 21
}
```



correct; unclear without  
looking at implementation

## Modular Verification

- Inspect method contracts
- Do *not* inspect method implementations
- *Design decision*

What are pros and  
cons of using modular  
verification?

# Verifying Method Calls

```
method triple(x: Int) returns (r: Int)
  requires x > 0
  ensures r > x
{
  r := 3 * x
}
```

```
method client() {
  var z: Int
  z := triple(7)
  assert z > 5
  assert z == 21
}
```



correct; unclear without  
looking at implementation

## Modular Verification

- Inspect method contracts
- Do *not* inspect method implementations
- *Design decision*

## Pros:

- Avoid client re-verification if implementation changes
- Respects the *information hiding* principle (encapsulation)
- Handling of recursion

## Cons:

- False negatives (*incompleteness*)
- Need to write more contracts

# Abstract Methods

```
method triple(x: Int) returns (r: Int)
  ensures r == 3 * x
```

```
method isqrt(x: Int) returns (r: Int)
  requires x >= 0
  ensures x >= r * r
  ensures x < (r+1) * (r+1)
```

```
method foo(a: Int) returns (b: Int)
  requires a > 0
  ensures b > a
{
  b := isqrt(a)
  b := triple(a)
}
```



- Contracts without Implementations
  - abstract from hard-to-verify code
  - abstract from unknown implementation
- Verification and good software engineering facilitate each other
  - *Incremental development* by refinement
  - Contracts become simpler if every method has a *single responsibility*
  - Avoid premature optimizations

## More abstract methods

```
method unsound(x: Int)
    returns (r: Int)
    ensures r != r
```

```
method test() {
    var a: Int
    a := unsound(17)
    assert 2 != 2
}
```

# More abstract methods

```
method unsound(x: Int)
  returns (r: Int)
  ensures r != r
```

```
method test() {
  var a: Int
  a := unsound(17)
  assert 2 != 2
}
```



- **Trusted code base:** code that is not checked by the verifier
- Danger of *unsoundness*: trusted inconsistencies may cause **false positives**
- Requires separate correctness arguments
- Methods are trusted until implemented

# Wrap-up: Informal Overview

- Specification mechanisms
  - Assertions
  - Pre- and postconditions
  - Underspecification
- Using an automated verifier
  - Modular reasoning with contracts
  - Abstract methods
  - Soundness and completeness issues
  - Trusted code base
- Verification and good software engineering facilitate each other
  - Information hiding, single responsibility principle
  - Incremental development