

Program Analysis for Software Security

Lecture 11

PERMISSION MODELS


Previously...

Heap data structures pose three major challenges for sequential verification

- Reasoning about aliasing
 - Permissions and separating conjunction 
- Framing, especially for dynamic data structures
 - Sound frame rule, predicates 
- Writing specifications that preserve information hiding
 - Data abstraction, heap-dependent functions 

Data abstraction

- To write implementation-independent specifications, we map the concrete data structure to mathematical concepts and specify the behavior in terms of those



- Mapping can be a function or a relation

Data abstraction via predicate arguments

```
method append(this: Ref, e: Int, c: Seq[Int])
  requires list(this, c)
  ensures list(this, c ++ Seq(e))
{
  unfold list(this, c)
  if(this.next == null) {
    var n: Ref
    n := new(*)
    n.next := null
    this.elem := e
    this.next := n
    fold list(n, Seq[Int]())
  } else {
    append(this.next, e, c[1..])
  }
  fold list(this, c ++ Seq(e))
}
```

- Solution supports information hiding
- No unfolding-expressions required in specifications
- Preconditions would ideally quantify existentially over abstract value
 - method append(this: Ref, e: Int)
 requires exists c: Seq[Int] ::
 list(this, c)
 ensures list(this, c ++ Seq(e))
- Since existentials are problematic for automation, ghost arguments are required instead (overhead!)

Reminder: partial functions

- Preconditions of heap-dependent functions specify the read effect
- Like method preconditions, they may also constrain the function arguments (including the heap)

```
function length(this: Ref): Int
  requires list(this)
{ ... }
```

```
function first(this: Ref): Int
  requires list(this) && 0 < length(this)
{
  content(this)[0]
}
```

Wrap-up: data abstraction

- Ghost fields + invariants
 - Manual updates of ghost state
 - Predicates cannot be abstract
- Predicate arguments
 - Support information hiding
 - Require ghost parameters
- Heap-dependent functions
 - Support information hiding
 - Typically exist in programs anyway (getters)
 - **Separation of concerns:**
predicates for permissions and invariants,
functions for abstraction


```
predicate list(this: Ref) {  
    acc(this.elem) && acc(this.next) &&  
    (this.next != null ==> list(this.next))  
}
```

```
function length(this: Ref): Int  
    requires list(this)  
{ ... }
```

```
function first(this: Ref): Int  
    requires list(this) && 0 < length(this)  
{ ... }
```

```
function content(this: Ref): Seq[Int]  
    requires list(this)  
{ ... }
```

Tentative course outline



Advanced permission models

1. Fractional permissions
2. Quantified permissions

Framing revisited

→ 00-clone.vpr

```
method cloneList(this: Ref) returns (res: Ref)
  requires list(this) // read only
  ensures list(this) && list(res)
  ensures content(this) == old(content(this))
{
  res := new(*)
  unfold list(this)
  if(this.next == null) {
    res.next := null
  } else {
    var tmp: Ref
    tmp := cloneList(this.next)
    res.elem := this.elem
    res.next := tmp
  }
  fold list(this)
  fold list(res)
}
```

- Methods that only read a data structure must specify that each abstraction remains **unchanged**
- Adding an abstraction requires changes to existing specifications (**non-modular**)
- Possible solution: specify that predicate version remains unchanged (not possible in Viper)
- We introduce a more expressive solution in the following

Fractional permissions

- To distinguish read and write access, permissions can be split and re-combined
- A permission amount π is a rational number in $[0,1]$
- Viper syntax
 - Permissions are fractions n/d
 - `write` for 1/1 and `none` for 0/1
 - `acc(E.f)` is a shortcut for `acc(E.f, write)`
 - `P(E)` is a shortcut for `acc(P(E), write)`
- Field read requires a non-zero permission
- Field write requires full (`write`) permission

Predicates (or assertions)

P ::= ...
| acc(E.f, π)
| acc(P(\bar{E}), π)

inhale acc(x.f, 1/2)
v := x.f



inhale acc(x.f, 1/2)
x.f := v



Manipulating fractional permissions

- Separating conjunction sums up permissions of the conjuncts

`acc(x.f, 1/2) && acc(x.f, 1/2)` is equivalent to `acc(x.f, 1/1)`

- inhale adds permissions

```
method cloneList(this: Ref) returns (res: Ref)
  requires acc(list(this), 1/2) // read only
  ensures acc(list(this), 1/2) && list(res)
{ ... }
```

- exhale *subtracts* permissions and
havocs only when *all* permission to
a location or predicate instance is
removed

```
method frameList(this: Ref) returns (l: Ref)
  requires list(this)
{
  var tmp1: Seq[Int]
  tmp1 := content(this)
  l := cloneList(this) // no havoc of version
  assert tmp1 == content(this)
}
```

- Values are framed as long as *some*
permission is held

Predicates and fractional permissions

→ 02-clone.vpr
→ 02-fractions.vpr

- Predicates may contain fractional permissions, e.g. to permit sharing
- Field locations with more than full permission are infeasible (magic)
- Predicate instances with more than full permission are feasible (no magic)
- Unfold and fold multiply the fraction of the predicate with the fractions in the predicate body

```
predicate readCell(this: Ref) {  
    acc(this.cell) && acc(this.cell.val, 1/2)  
}
```

```
predicate P(this: Ref) {  
    acc(this.val, 1/2)  
}
```

```
inhale acc(x.val)  
fold P(x)  
fold P(x)  
exhale P(x) && P(x) // not false
```



```
inhale acc(readCell(x), 1/4)  
unfold acc(readCell(x), 1/4)  
exhale acc(x.cell.val, 1/8)
```



Heap-dependent functions

- Heap-dependent functions may only read the heap
- Hence, an arbitrarily small fraction would be sufficient
- Problem: we don't know how often permissions are split
- Possible solution: use **wildcard** to avoid concrete fraction

```
function length(this: Ref): Int
    requires list(this)
{
    unfolding list(this) in
        (this.next == null ? 0 : length(this.next) + 1)
}
```

```
inhale acc(list(this), 1/2)
x := length(this)
```



```
function length(this: Ref): Int
    requires acc(list(this), wildcard)
{
    unfolding acc(list(this), wildcard) in
        (this.next == null ? 0 : length(this.next) + 1)
}
```

Adjusted encoding: permissions and field access

- Permissions are tracked in a global permission mask

```
type MaskType = Map<T>[(Ref, Field T), Real]  
var Mask: MaskType
```

- Convention: $\text{Mask}[\text{null}, f] == 0.0$ for all fields f

- Field access

```
v := x.f
```

```
assert Mask[x,f] > 0.0  
v := Heap[x,f]
```

```
x.f := E
```

```
assert Mask[x,f] == 1.0  
Heap[x,f] := E
```

- Field access requires permission!

Adjusted encoding: inhale

- **inhale A** means:
 - obtain all permissions required by assertion A
 - assume all logical constraints
- Encoding is defined recursively over the structure of A

inhale acc(E.f, π)

$\text{Mask}[[[E]], f] := \text{Mask}[[[E]], f] + \pi$
assume $\text{Mask}[[[E]], f] \leq 1.0$

Reaching more than full permission
for a **field** location goes to magic

inhale acc(P(E), π)

$\text{Mask}[\text{null}, \text{PField}([[E]])] := \text{Mask}[\text{null}, \text{PField}([[E]])] + \pi$

inhale A && B

[[inhale A]]; [[inhale B]]

Separating conjunction:
add sum of permissions

- The encoding also asserts that E and π are well-defined (omitted here)

Adjusted encoding: exhale

- **exhale A** means:
 - assert all logical constraints
 - check and remove all permissions required by assertion A
 - havoc any locations to which all permission is lost
- Encoding is defined recursively over the structure of A

exhale acc(E.f, π)

```
assert Mask[[[E]],f] >=  $\pi$ 
Mask[[[E]],f] := Mask[[[E]],f] -  $\pi$ 
```


exhale A

```
var oldMask: MaskType
var newHeap: HeapType
oldMask := Mask
[[exhale A]]
assume forall y,g :: Mask[y,g] > 0.0 ==> newHeap[y,g] == Heap[y,g]
Heap := newHeap // effectively havocs all Locations to which all
                permission was lost
```

Sharing in data structures


- Full permissions can describe tree-shaped data structures only

```
predicate exclusiveCell(this: Ref) {  
    acc(this.cell) && acc(this.cell.val)  
}
```



- Fractional permissions allow sharing

```
predicate sharedCell(this: Ref) {  
    acc(this.cell) && acc(this.cell.val, 1/2)  
}
```




- Sharing is used in many data structures
 - Examples: doubly-linked lists, global data, caches, graphs, ...

Case study: binary reference counting

- Binary reference counting optimizes code that uses immutable data
- Keep track whether the data is shared
- Updates on shared data perform a copy (“copy on write”)
- Updates on unshared data perform a destructive update
- Once shared, the data does not go back to unshared (unlike with full reference counting)

- Example: text segments



The Perm type


- The demo uses a ghost field of type `Perm`
- Values of type `Perm` include:
 - constants `none`, `write`, `wildcard`, and fractions
 - expressions, e.g., `write - x.frac` or `2*write`
- `Perm` is typically used for ghost variables
 - Parameterize methods that require read permission
 - Perform permission accounting when permissions are distributed and later re-collected
- Type `Perm` is encoded as a real

`field frac: Perm`

```
method cloneList(this: Ref, p: Perm)
           returns (res: Ref)
requires acc(list(this), p)
ensures acc(list(this), p) &&
list(res)
```

Sharing in mutable data structures

- Previously: immutable shared objects
- To specify mutable data structures with sharing, we arrange fractional permissions such that they can be combined to obtain a full permission



- Example: doubly-linked list

```
predicate nodes(this: Ref) {  
    acc(this.next) && acc(this.prev, 1/2) &&  
    (this.next != null ==>  
        acc(this.next.prev, 1/2) &&  
        this.next.prev == this &&  
        nodes(this.next)  
    )  
}
```


```
predicate dlist(this: Ref) {  
    acc(this.head) &&  
    (this.head != null ==>  
        acc(this.head.prev, 1/2) &&  
        this.head.prev == null &&  
        nodes(this.head)  
    )  
}
```

Advanced permission models


1. Fractional permissions
2. Quantified permissions

Limitations of recursive predicates


- Recursive predicates allow one to specify unbounded data structures
 - Traversals happen [in the order](#) in which the predicate needs to be unfolded
- Predicates are not ideal for many other use cases




Iterative traversals



Other traversal orders



Random-access data structures



Arbitrary cyclic data structures

Quantified permissions

- To denote permission to an unbounded set of locations without prescribing a traversal order, we allow permissions and predicates in universal quantifiers

Permissions

$P ::= \dots \mid \text{forall } \overline{x:T} :: P$

- Universal quantifiers can be thought of as a possibly-infinite iterated conjunction


$\text{forall } x:T :: P \iff P[x/v_1] \wedge P[x/v_2] \wedge \dots$

- Viper's `forall` represents a possibly-infinite **iterated separating conjunction**

$\text{forall } x:T :: P \iff P[x/v_1] * P[x/v_2] * \dots$

Explicit footprints

- As alternative to predicates, we can specify permission to an unbounded set of locations by
 - maintaining an explicit set of references as ghost state (the explicit footprint)
 - quantifying over the set elements in specifications



```
field head: Ref
field nodes: Set[Ref] // explicit footprint
```

```
predicate list(this: Ref) {
  acc(this.head) && acc(this.nodes) &&
  (forall n: Ref :: n in this.nodes ==> acc(n.elem) && acc(n.next) &&
   (n.next != null ==> n.next in this.nodes)) &&
  (this.head != null ==> this.head in this.nodes)
}
```

Well-formed quantified permissions

- Viper requires for each assertion $\text{acc}(E.f)$ under a **forall** $x:T$ that E is **injective** for all instantiations x_1, x_2 :
$$x_1 \neq x_2 \implies E[x/x_1] \neq E[x/x_2]$$
- Analogous rule applies to predicates (for parameter tuples)
- Examples

forall $x: \text{Ref} :: x \in s \implies \text{acc}(x.f) \quad // \ s \ has \ type \ \text{Set}[\text{Ref}]$



forall $x: \text{Ref} :: \text{acc}(y.f) \ \&& \ (y.f \neq x \implies P(x))$




forall $i: \text{Int} :: 0 \leq i \ \&& \ i < |t| \implies \text{acc}(t[i].f) \quad // \ t \ has \ type \ \text{Seq}[\text{Ref}]$

well-formed if t contains no duplicates

Complex sharing: graph marking


- We represent a graph as a set of nodes
- Each node stores a (possibly empty) set of successors
- Each node contains a flag that is set during marking




```
field next: Set[Ref]
field flag: Bool

define graph(nodes) (
  forall n: Ref :: n in nodes ==> acc(n.next) && acc(n.flag) && (n.next subset nodes)
)
```


Quantified permissions address the limitations of predicates




Iterative traversals



Other traversal orders



Random-access data structures



Arbitrary cyclic data structures

Arrays

- Viper does not have built-in arrays
- In contrast to sequences, arrays are mutable heap data structures
- We model arrays by a set of disjoint references that can be accessed via an index
- `loc(a, i).val` models `a[i]`
- More-dimensional arrays can be encoded analogously

```
field val: Int // for integer arrays

domain Array {
    function loc(a: Array, i: Int): Ref
    function len(a: Array): Int
    function first(r: Ref): Array
    function second(r: Ref): Int
}

axiom injectivity {
    forall a: Array, i: Int :: {loc(a, i)}
        first(loc(a, i)) == a &&
        second(loc(a, i)) == i
}

axiom length_nonneg {
    forall a: Array :: len(a) >= 0
}
```

Accessing array locations

→ 09-array-inc-all.vpr

- Arrays are random-access data structures
- We can express permissions using quantified permissions

```
forall i: Int :: 0 <= i && i < len(a) ==> acc(loc(a, i).val)
```

- Similarly for sub-ranges of the array

- We define macros for convenient access

```
define lookup(a, i)
    loc(a, i).val
```

```
define update(a, i, e) {
    loc(a, i).val := e
}
```

- Bounds are checked implicitly via permissions

Wrap-up: advanced permission models

- Fractional permissions
 - Distinguish between read and write permission
 - Are useful to express sharing, to strengthen framing, and for concurrency (see later)
- Quantified permissions
 - Complement predicates for the specification of unbounded data structures
 - Are especially useful for random-access structures, complex sharing, and flexible traversals
 - Inherit challenges of quantification (controlling instantiations, performance)
- Other permission models exist
 - Magic wands (permission-aware implication): useful to specify partial data structures
 - Counting permissions are related to fractional permissions, but use units